HFENet: Hierarchical Feature Alignment for Data-Efficient Fine-Grained Image Classification

Shuo Wu[1,2*], Hongjie Huang[1], Zhonghao Li[1], Chenlong Wang[1], Qijuan Gao[1], Yu Su[1,2]

[1*]School of Computer Science and Artificial Intelligence, Hefei Normal University, Hefei, China.

[2]Key Laboratory of Philosophy and Social Science of Anhui Province on Adolescent Mental Health and Crisis Intelligence Intervention, China.

*Corresponding author(s). E-mail(s): wushuo@ustc.edu.cn;
Contributing authors: 2242128115@qq.com; anthony23six@gmail.com;
1264910799@qq.com; gaoqijuan@hotmail.com; yusu@hfnu.edu.cn;

Abstract

Fine-grained image classification aims to distinguish visually similar subcategories within the same basic category and remains challenging under limited data and low-resolution inputs, where reliable discriminative representations are difficult to learn. Existing methods have introduced hierarchical semantic cues, such as coarse-grained labels, but most rely on simple cascaded, parallel, or feature-fusion designs without explicitly modeling cross-level semantic dependency or suppressing feature redundancy across branches. To address these issues, we propose HFENet for fine-grained image classification. The core idea is hierarchical feature alignment through cross-branch interaction between coarse-grained and fine-grained representations. Specifically, a cross-branch attention mechanism transfers semantic guidance from the coarse branch to the fine branch, while an orthogonality regularization term reduces redundant representations and promotes complementary feature learning. In addition, a hierarchical Kullback–Leibler divergence is incorporated to improve semantic consistency between branches during training. Experiments on CIFAR-100 show that HFENet achieves accuracies of 70.8%, 77.2%, 80.5%, and 82.7% under 30%, 50%, 80%, and 100% training ratios, respectively. These results demonstrate consistent improvements over the baseline dual-branch model and representative comparison methods, especially in data-constrained settings. Ablation studies and visualization analyses further verify the effectiveness of hierarchical alignment and redundancy suppression.

Keywords: Fine-grained image classification, Hierarchical feature alignment, Cross-branch attention, Orthogonality regularization, Complementary representation learning, Data-efficient learning

1 Introduction

Fine-grained image classification (FGIC) aims to distinguish visually similar subcategories within the same basic-level category. Compared with generic image classification, FGIC is more challenging due to subtle inter-class differences and large intra-class variations. These difficulties become more severe under limited training data and low-resolution inputs, where discriminative cues

are often weak and models are more prone to overfitting.

Existing studies have improved FGIC by introducing attention mechanisms, multi-branch architectures, and hierarchical semantic supervision. However, two limitations remain. First, the relationship between coarse-grained and fine-grained representations is often not explicitly modeled, since many hierarchical methods simply concatenate or sum multi-level features without capturing the conditional dependency that coarse semantics should guide fine-grained discrimination. Second, multi-branch designs may introduce redundant representations, as different branches tend to learn overlapping information without explicit constraints, resulting in suboptimal feature utilization. Recent studies have further advanced fine-grained recognition through refined attention design, semantic consistency modeling, and structured hierarchical reasoning [17–19].

To address these issues, we propose HFENet for fine-grained image classification. The framework is built around a central idea: hierarchical feature alignment for coarse-to-fine representation learning under constrained settings. Specifically, a cross-branch attention (CBA) module transfers coarse-level semantic guidance to the fine-grained branch through conditional feature modulation, while an orthogonality regularization term suppresses cross-branch redundancy and promotes complementary representation learning. In addition, a hierarchical KL consistency objective is introduced to improve semantic coherence between coarse- and fine-grained predictions during training. Together, these components enable the network to learn more structured and data-efficient representations, especially in limited-data and low-resolution scenarios.

The main contributions of this work are summarized as follows:

- We propose HFENet, a dual-branch framework for fine-grained image classification under limited-data and low-resolution settings. The core idea is to achieve hierarchical feature alignment between coarse-grained and fine-grained representations.
- We design a cross-branch attention mechanism that transfers semantic guidance from the coarse branch to the fine branch through joint channel-spatial modulation, enabling adaptive coarse-to-fine conditioning beyond conventional intra-branch attention.
- We introduce an orthogonality regularization term to suppress cross-branch redundancy and encourage complementary feature learning, thereby improving the discriminative efficiency of the dual-branch architecture.
- We incorporate hierarchical KL consistency into the joint optimization objective and conduct experiments, ablation studies, and visualization analyses to validate the proposed framework.

2 Related Work

2.1 Fine-Grained Image Classification

Fine-grained image classification (FGIC) aims to distinguish subordinate categories with subtle visual differences. Early studies mainly relied on part localization and region-based modeling to capture discriminative local cues, such as Part-based R-CNN [8] and recurrent attention-based methods [6]. Subsequent work explored stronger feature interaction mechanisms, including bilinear pooling [9], multi-attention learning [5], and progressive multi-granularity training [1]. Transformer-based methods such as TransFG [7] further improved representation learning by modeling long-range dependencies. Recent studies have continued to strengthen fine-grained representation learning through refined attention design, semantic consistency enhancement, and higher-order feature interaction, such as CRA-SMRA [17], CFGVC-Net [18], and high-order feature interaction modeling [22]. Despite these advances, many FGIC approaches still mainly improve discrimination within a single representation stream and do not explicitly model coarse-to-fine semantic conditioning across hierarchical levels.

2.2 Attention Mechanisms in Visual Recognition

Attention mechanisms have been widely adopted to improve visual feature learning. Representative examples include channel attention in SE-Net [3], joint channel-spatial attention in CBAM [4], and long-range dependency modeling in non-local

networks [10]. These methods mainly recalibrate features within the same branch or stage, and are therefore well suited for intra-stream feature enhancement. However, in hierarchical recognition problems, the key requirement is not only to highlight salient responses, but also to explicitly transfer semantic guidance across levels. This motivates the use of a cross-branch attention mechanism in our framework.

2.3 Hierarchical Learning in Image Classification

Hierarchical learning has been explored to exploit semantic structure in image recognition. HD-CNN [2] uses a cascaded hierarchy to progressively refine predictions, while PMG [1] introduces progressive granularity learning to strengthen fine-grained discrimination. More recent studies have revisited hierarchical supervision from the perspective of cross-level consistency and multi-granularity reasoning. For example, CHBC [19] improves fine-grained visual classification by enforcing bidirectional consistency across different hierarchy levels, and recent exploratory work on free-grained hierarchical recognition also suggests the practical relevance of learning under mixed coarse- and fine-grained supervision [21]. These studies show that hierarchical supervision can improve recognition performance. Nevertheless, existing methods often rely on cascaded routing, consistency constraints, or implicit multi-level learning, rather than explicitly modeling directional semantic interaction between parallel coarse and fine branches. Moreover, redundancy across branches is rarely addressed as a first-class design objective.

2.4 Transformer and Modern CNN Architectures

Recent vision backbones such as ViT [11], Swin Transformer [12], PVTv2 [15], ConvNeXt [14], and ConvNeXt V2 [16] have achieved strong performance on large-scale benchmarks. However, these architectures often benefit substantially from large-scale pre-training or extensive data, and their advantages may be less pronounced in limited-data settings when trained from scratch [13]. In this work, we compare HFENet with representative Transformer and modern CNN backbones under a unified from-scratch protocol, with the goal of evaluating relative data efficiency rather than claiming universal superiority across model families.

2.5 Comparison with Existing Methods

Compared with existing FGIC, attention-based, and hierarchical classification methods, HFENet is built around a more explicit design objective: hierarchical feature alignment for coarse-to-fine representation learning. Unlike standard attention modules that operate within a single branch, our cross-branch attention uses coarse-level semantics to condition fine-grained feature modulation. Unlike conventional dual-branch designs that rely on implicit specialization, our framework further introduces orthogonality regularization to suppress redundant cross-branch representations. These differences can be understood from the perspectives of hierarchical supervision, cross-level interaction, and redundancy suppression. Recent studies have also emphasized fine-grained alignment and structured semantic reasoning, for example by modeling fine-grained visual-concept relations or hierarchical consistency [19, 20]. In contrast, HFENet focuses on coarse-to-fine feature alignment in a dual-branch architecture and explicitly combines directional cross-branch modulation with redundancy suppression.

3 Method

3.1 Design Motivation and Representation Objectives

Fine-grained image classification can be viewed as a hierarchical representation learning problem. Coarse-grained categories provide high-level semantic structure, whereas fine-grained categories require local discriminative cues. Given an input image $\mathbf{X}$, its fine-grained label $y_f \in \{1, \ldots, C_f\}$ and coarse-grained label $y_c \in \{1, \ldots, C_c\}$ with $C_f > C_c$, the desired representation properties of the two branches can be formulated as:

$$\begin{aligned} &I(\mathbf{F}_c; y_c) \gg I(\mathbf{F}_c; y_f), \\ &I(\mathbf{F}_f; y_f) \gg I(\mathbf{F}_f; y_c), \\ &I(\mathbf{F}_c; \mathbf{F}_f) \approx 0, \end{aligned} \tag{1}$$

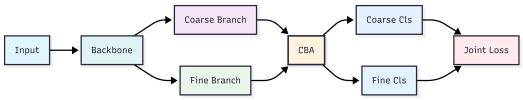

Fig. 1: Overview of HFENet, including the shared backbone, the coarse branch, the fine branch, and the associated training objectives

where $I(\cdot;\cdot)$ denotes mutual information. These expressions should be understood as representation objectives rather than explicitly optimized information-theoretic quantities. They indicate that the coarse branch is expected to encode category-level semantic structure, the fine branch should emphasize subtle discriminative details, and the two branches should remain complementary rather than redundant.

Based on these objectives, HFENet is designed around two key ideas: hierarchical alignment and redundancy suppression. The former is implemented through cross-branch attention, which allows coarse semantics to condition fine-grained feature modulation, while the latter is achieved by orthogonality regularization, which discourages the two branches from learning duplicated representations.

3.2 Overall Architecture

The proposed HFENet is built around a central objective: hierarchical feature alignment for coarse-to-fine representation learning. Given shared backbone features, the network decomposes representation learning into a coarse branch and a fine branch, so that high-level category semantics and fine-grained discriminative cues can be modeled in a coordinated manner. To realize this objective, HFENet incorporates two core mechanisms and one auxiliary optimization strategy. The first mechanism is cross-branch attention (CBA), which enables directional semantic guidance from the coarse branch to the fine branch. The second mechanism is orthogonality regularization, which suppresses redundant cross-branch representations and encourages complementarity. In addition, hierarchical KL consistency is used during training to improve prediction-level semantic coherence. The overall architecture is shown in Fig. 1.

3.3 Cross-Branch Attention Mechanism

3.3.1 Design Motivation

Conventional attention modules such as SE and CBAM mainly recalibrate features within the same branch. Although effective for enhancing local saliency, they do not explicitly model the directional dependency that coarse-grained semantics should constrain fine-grained discrimination in hierarchical recognition tasks. Simple feature concatenation or linear fusion is also insufficient, because it lacks adaptive conditioning and may introduce redundant information sharing across branches.

To address this issue, we design a cross-branch attention mechanism that uses coarse-level semantics as conditional guidance for fine-grained feature modulation. From an information conditioning perspective, this mechanism can be interpreted as increasing the discriminative utility of fine-grained representations under the guidance of coarse semantics:

$$I(\mathbf{F}'_f; y_f \mid \mathbf{F}_c), \tag{2}$$

where $\mathbf{F}'_f$ denotes the modulated fine-grained features.

3.3.2 Mechanism Implementation

To enable conditional semantic guidance, we design a cross-branch attention module that uses coarse features to modulate fine features. The coarse feature $\mathbf{F}_c$ is compressed into a channel attention vector $\mathbf{a}_c \in \mathbb{R}^{1\times1\times C}$ and a spatial attention map $\mathbf{A}_s \in \mathbb{R}^{H\times W\times1}$. By broadcast multiplication, they are combined into a channel-spatial modulation mask

$$\mathbf{M} = 1 + \alpha \cdot (\mathbf{A}_s \odot \mathbf{a}_c), \tag{3}$$

where $\mathbf{M} \in \mathbb{R}^{H\times W\times C}$ and α is a learnable scalar. The fine-grained feature is then modulated as

$$\mathbf{F}'_f = \mathbf{F}_f \odot \mathbf{M}. \tag{4}$$

This process implements spatially- and channel-wise adaptive guidance from the coarse branch to the fine branch.

3.3.3 Comparison with Existing Attention Mechanisms

To highlight the distinctiveness of our cross-branch attention, Table 1 compares it with several representative attention mechanisms.

As summarized in Table 1, existing attention mechanisms mainly operate within a single branch or do not explicitly encode coarse-to-fine hierarchical conditioning. In contrast, the proposed CBA is designed for cross-branch modulation, where coarse-level semantics provide adaptive channel-spatial guidance to fine-grained representations.

3.4 Orthogonality Regularization

A dual-branch architecture does not automatically guarantee complementary representation learning. When both branches are optimized on the same input and related supervision signals, they may converge to overlapping feature subspaces, which weakens the benefit of hierarchical decomposition. Therefore, redundancy suppression is not merely an optional regularization term, but a necessary constraint for preserving functional differentiation between the coarse and fine branches.

To reduce cross-branch redundancy, we first project the coarse and fine features into a shared embedding space. For a mini-batch of N samples, let $\mathbf{F}_c \in \mathbb{R}^{N\times d_c}$ and $\mathbf{F}_f \in \mathbb{R}^{N\times d_f}$ denote the batch-wise coarse and fine features, respectively. We introduce two linear projections $\phi_c(\cdot)$ and $\phi_f(\cdot)$ to obtain

$$\hat{\mathbf{F}}_c = \phi_c(\mathbf{F}_c), \qquad \hat{\mathbf{F}}_f = \phi_f(\mathbf{F}_f), \tag{5}$$

where $\hat{\mathbf{F}}_c, \hat{\mathbf{F}}_f \in \mathbb{R}^{N\times d}$ are mapped to the same feature dimension d.

The cross-branch correlation matrix is then defined as

$$\mathbf{C}_{cf} = \frac{1}{N}\hat{\mathbf{F}}_c^{\top}\hat{\mathbf{F}}_f. \tag{6}$$

The orthogonality regularization is formulated as

$$\mathcal{L}_{\text{orth}} = \|\mathbf{C}_{cf}\|_F^2. \tag{7}$$

Minimizing Eq. (7) suppresses correlated responses between the two branches and encourages them to encode complementary rather than redundant information. This constraint is particularly important in dual-branch architectures with shared low-level features, where branch collapse toward similar representations can otherwise occur.

3.5 Hierarchical Joint Optimization

Let $\mathbf{z}_f \in \mathbb{R}^{C_f}$ and $\mathbf{z}_c \in \mathbb{R}^{C_c}$ denote the logits of the fine and coarse classifiers, respectively. Their posterior probabilities are defined as

$$\mathbf{p}_f = \text{softmax}(\mathbf{z}_f), \qquad \mathbf{p}_c = \text{softmax}(\mathbf{z}_c). \tag{8}$$

Because each fine-grained class belongs to exactly one coarse-grained category, the fine prediction can be aggregated into the coarse label space. Let $\mathcal{G}(k)$ denote the set of fine classes belonging to coarse class k. The induced coarse distribution from the fine classifier is defined as

$$\tilde{p}_c(k) = \sum_{j\in\mathcal{G}(k)} p_f(j), \qquad k = 1, \ldots, C_c. \tag{9}$$

Based on this mapping, the hierarchical KL consistency loss is formulated as

$$\mathcal{L}_{\text{KL}} = \sum_{k=1}^{C_c} \tilde{p}_c(k) \log \frac{\tilde{p}_c(k)}{p_c(k) + \varepsilon}, \tag{10}$$

where ε is a small constant for numerical stability. This term encourages semantic consistency between the fine-grained prediction and the coarse-grained prediction, thereby improving cross-level coherence during training.

The overall loss function combines fine-grained classification loss $\mathcal{L}_f$, coarse-grained classification loss $\mathcal{L}_c$, hierarchical KL consistency loss $\mathcal{L}_{\text{KL}}$, and orthogonality regularization:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_f + \lambda_c\mathcal{L}_c + \lambda_{\text{KL}}\mathcal{L}_{\text{KL}} + \lambda_{\text{orth}}\mathcal{L}_{\text{orth}}, \tag{11}$$

where λ_c, λ_{KL}, and λ_{orth} are balancing hyperparameters.

The model is trained end-to-end using the overall objective in Eq. (11).

4 Experiments

4.1 Datasets and Evaluation Metrics

We evaluate the proposed method on CIFAR-100 and Stanford Cars. CIFAR-100 contains 100 fine-grained classes, which are grouped into 20

Table 1: Comparison of CBA with representative attention mechanisms

Method	Input	Target	Hier.	Dyn.	Cost
SE-Net	Single	Same	No	Static	$O(c^2/r)$
CBAM	Single	Same	No	Static	$O(c^2/r + hwc)$
Non-local	Single	Same	No	Dynamic	$O((hw)^2c)$
Cross-attn.	Two-branch	Other	Weak	Dynamic	$O(hwc^2)$
CBA (Ours)	**Cross-level**	**Fine**	**Strong**	**Dynamic**	$\mathbf{O(c^2/r + hwc)}$

coarse-grained superclasses and are used to study performance under different training ratios (30%, 50%, 80%, and 100%) in a low-resolution setting. Stanford Cars contains 196 fine-grained classes. For hierarchical supervision, we derive 49 coarse-grained categories based on vehicle brands from the original class annotations. Following common practice, Top-1 classification accuracy is used as the evaluation metric for all experiments.

4.2 Implementation Details

All models are implemented in PyTorch and trained from scratch under a unified setting. WideResNet-16-8 is used as the backbone for HFENet and the main baseline variants. The optimizer is AdamW with an initial learning rate of 0.001 and weight decay of 5×10^{-4}, and cosine annealing is applied for 300 epochs. The batch size is set to 64. Data augmentation includes random cropping, horizontal flipping, AutoAugment, and random erasing ($p = 0.5$). The loss weights are set to $\lambda_c = 0.5$, $\lambda_{\mathrm{KL}} = 0.2$, and $\lambda_{\mathrm{orth}} = 0.0005$, determined through grid search.

To keep the evaluation consistent with the target problem setting, CIFAR-100 is trained and tested at its original low resolution of 32×32, while Stanford Cars is resized to 128×128. This design is intended to evaluate the proposed method under constrained visual conditions, rather than to pursue the highest possible accuracy under unrestricted high-resolution settings.

4.3 Comparison under a Unified Training Protocol

All comparison models are trained under a unified from-scratch protocol to evaluate relative data efficiency under limited-data and low-resolution conditions. Under this setting, HFENet consistently outperforms the baseline WRN-16-8 and the compared representative backbones across different training ratios. The gains are particularly pronounced when supervision is limited, suggesting that hierarchical feature alignment provides a useful inductive bias when fine-grained cues are insufficient or weakened by constrained visual resolution. It should be emphasized that Transformer and modern CNN backbones are evaluated here without large-scale pre-training. Therefore, the results are mainly intended to compare data efficiency and representation robustness under a unified constrained training regime, rather than to establish a universal ranking across different model families or unrestricted high-resolution settings.

Results in Table 2 show that HFENet consistently improves over the baseline WRN-16-8 and the dual-branch variant without interaction across all training ratios. The improvement is particularly significant in limited-data settings, where HFENet exceeds WRN-16-8 by 12.4 points at 30% data and 8.5 points at 50% data. The margin becomes smaller at 100% data, which suggests that the benefit of hierarchical alignment is most pronounced when supervision is limited and discriminative cues are relatively difficult to stabilize. Notably, HFENet at 50% data already surpasses the full-data performance of WRN-16-8, further indicating its advantage in data-efficient fine-grained recognition under constrained conditions.

4.3.1 Convergence Analysis

Convergence curves in Fig. 2 further illustrate the optimization behavior. HFENet achieves faster initial convergence and better final accuracy, demonstrating that hierarchical guidance and redundancy suppression can improve optimization stability in limited-data and low-resolution fine-grained recognition.

Table 2: Classification accuracy (%) on CIFAR-100 under different training ratios

Method	30% data	50% data	80% data	100% data
WRN-16-8 (Baseline)	58.4 ± 0.3	68.7 ± 0.4	76.3 ± 0.2	80.2 ± 0.2
Dual-branch WRN (no interaction)	61.6 ± 0.4	70.7 ± 0.3	78.2 ± 0.3	80.6 ± 0.2
ResNet50	55.2 ± 0.5	64.8 ± 0.4	72.5 ± 0.3	76.8 ± 0.3
DenseNet121	57.6 ± 0.4	66.9 ± 0.3	74.1 ± 0.3	78.3 ± 0.2
EfficientNet-B0	56.8 ± 0.5	66.2 ± 0.4	73.7 ± 0.3	77.9 ± 0.3
ViT-B	39.5 ± 0.7	49.4 ± 0.6	64.6 ± 0.4	69.7 ± 0.4
Swin-T	52.7 ± 0.6	59.7 ± 0.5	66.2 ± 0.4	69.1 ± 0.3
PVTv2	48.3 ± 0.6	56.9 ± 0.5	63.8 ± 0.4	67.5 ± 0.4
ConvNeXt V2	53.8 ± 0.5	62.4 ± 0.4	69.1 ± 0.3	72.6 ± 0.3
HFENet (Ours)	**70.8 ± 0.3**	**77.2 ± 0.2**	**80.5 ± 0.2**	**82.7 ± 0.2**

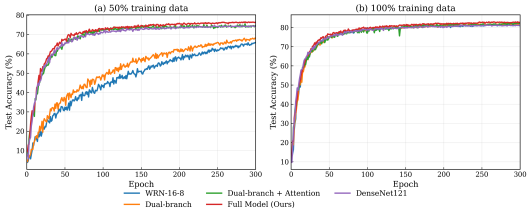

Fig. 2: Test accuracy convergence curves on CIFAR-100 with 50% and 100% training data

4.4 Results on Stanford Cars

We further evaluate HFENet on Stanford Cars, a standard fine-grained benchmark with 196 model classes. For this dataset, we derive 49 brand-level coarse categories to enable hierarchical supervision.

Table 3: Classification accuracy (%) on Stanford Cars

Method	50% data	100% data
WRN-16-8 (Baseline)	50.52 ± 0.6	81.17 ± 0.3
Dual-branch WRN (no interaction)	62.48 ± 0.5	82.24 ± 0.3
Dual-branch + CBA	63.25 ± 0.4	83.04 ± 0.2
HFENet (Ours)	**65.73 ± 0.4**	**83.95 ± 0.2**
ResNet50	42.07 ± 0.7	80.06 ± 0.4
DenseNet121	49.12 ± 0.6	82.31 ± 0.3

HFENet achieves the best performance among the compared methods on Stanford Cars in both training settings. Compared with WRN-16-8, the proposed method improves accuracy by 15.21 points at 50% data and by 2.78 points at 100% data, indicating that hierarchical feature alignment remains effective on a standard FGVC benchmark under reduced-resolution conditions. The larger gain at 50% data further suggests that the proposed framework is particularly beneficial when both supervision and fine-grained visual detail are constrained.

4.5 Ablation Studies

We conduct ablations on CIFAR-100 to isolate the contribution of each component. Results are presented in Table 4.

Dual-branch decomposition alone yields limited improvements over the baseline, indicating that merely separating coarse- and fine-level pathways is insufficient. Introducing CBA brings substantial gains, which confirms the importance of explicit coarse-to-fine semantic guidance. Adding orthogonality regularization further improves performance by suppressing redundant cross-branch representations. The full model achieves the best results across all training ratios, suggesting that hierarchical alignment and redundancy suppression are complementary rather than interchangeable.

4.6 Feature Correlation Analysis

To examine whether orthogonality regularization effectively reduces branch-wise redundancy, we visualize the correlation patterns of dual-branch features with and without the proposed regularization. As shown in Fig. 3, the model without orthogonality regularization exhibits stronger cross-branch correlation, indicating that the two branches tend to encode overlapping information. After introducing orthogonality regularization, the overall correlation magnitude is reduced, which suggests that the coarse and fine branches become more complementary.

This observation is consistent with the ablation results in Table 4, where adding orthogonality regularization yields further gains over the plain dual-branch architecture. These results support the claim that redundancy suppression is beneficial for preserving functional differentiation between the

Table 4: Ablation results on CIFAR-100 under different training ratios (% accuracy)

Ratio	WRN-16-8	Dual	Dual+CBA	Dual+Orth.	Full
30%	58.4 ± 0.3	61.6 ± 0.4	69.9 ± 0.3	62.3 ± 0.2	**70.8 ± 0.3**
50%	68.7 ± 0.4	70.7 ± 0.3	75.3 ± 0.3	71.2 ± 0.3	**77.2 ± 0.2**
80%	76.3 ± 0.2	78.2 ± 0.3	79.2 ± 0.2	78.8 ± 0.2	**80.5 ± 0.2**
100%	80.2 ± 0.2	80.6 ± 0.2	82.1 ± 0.2	81.2 ± 0.2	**82.7 ± 0.2**

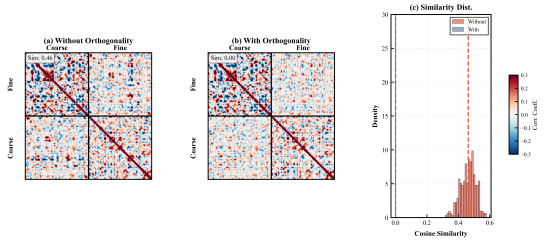

Fig. 3: Feature correlation analysis of dual-branch representations with and without orthogonality regularization. The regularized model exhibits weaker cross-branch correlation, indicating reduced redundancy and stronger complementarity

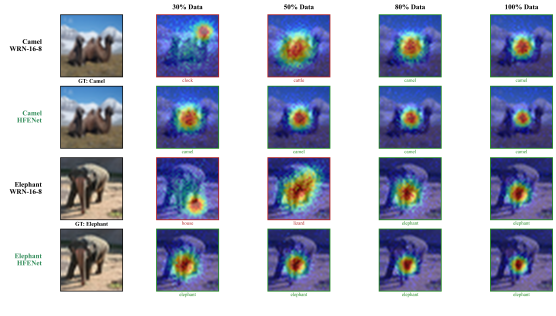

Fig. 4: Grad-CAM comparison between WRN-16-8 and HFENet on CIFAR-100 under different training ratios

two branches, especially when limited supervision and reduced visual detail make complementary representation learning more important.

4.7 Visualization Analysis

We visualize Grad-CAM heatmaps on CIFAR-100 (Fig. 4) to qualitatively examine the effect of hierarchical feature alignment. Compared with WRN-16-8, HFENet tends to produce more target-focused responses and is less easily distracted by background regions or irrelevant structures.

On CIFAR-100, the attention maps of HFENet are generally more concentrated on the object body and its discriminative parts, whereas the baseline model often produces scattered or background-sensitive responses. This difference is more evident under reduced training ratios, which is consistent with the quantitative results showing stronger gains in limited-data settings. These visualizations suggest that coarse-to-fine semantic guidance helps stabilize discriminative localization when fine-grained supervision is insufficient and local visual evidence is weakened by low-resolution inputs.

4.8 Error Analysis

We further inspect representative failure cases of HFENet on CIFAR-100 to better understand its remaining limitations. Most misclassifications can be grouped into three categories. First, severe occlusion of discriminative parts prevents the fine branch from capturing reliable local evidence. Second, some categories exhibit extremely subtle inter-class differences, for which coarse-level guidance alone is insufficient to resolve ambiguity. Third, low-resolution samples limit the availability of local structural details and reduce the effectiveness of fine-grained discrimination.

These observations indicate that, although hierarchical guidance improves branch coordination and discriminative focus, the model still depends on sufficiently visible local cues. They also suggest that future improvement may require stronger part-aware modeling, higher-resolution feature extraction, or more robust local detail enhancement.

5 Conclusion

This paper proposed HFENet, a dual-branch framework for fine-grained image classification built around the central idea of hierarchical feature alignment for coarse-to-fine representation learning. Within this framework, cross-branch attention enables directional semantic guidance from the coarse branch to the fine branch, while orthogonality regularization suppresses redundant cross-branch representations and promotes complementarity. Hierarchical KL consistency further improves semantic coherence during training. Experimental results on CIFAR-100 and

Stanford Cars show that HFENet consistently improves recognition performance, especially under limited-data and reduced-resolution settings. These findings suggest that explicitly combining hierarchical alignment and redundancy suppression is a promising direction for data-efficient and robust fine-grained recognition under constrained visual conditions.

Declarations

Funding

This work is a research outcome of the Major Project of Philosophy and Social Science Planning in Anhui Province, China (Project Approval No. AHSKZD2024D03).

Competing interests

The authors declare no competing interests.

Data availability

CIFAR-100 and Stanford Cars are publicly available from their official sources. Additional experimental results are available from the corresponding author upon reasonable request.

Code availability

Code is available from the corresponding author upon reasonable request.

Author contributions

Shuo Wu contributed to conceptualization, supervision, and manuscript revision. Hongjie Huang contributed to methodology, experiments, data analysis, and manuscript drafting. Zhonghao Li contributed to experimental validation and result analysis. Chenlong Wang contributed to implementation and visualization. Qijuan Gao contributed to manuscript review and editing. Yu Su contributed to supervision and project administration. All authors read and approved the final manuscript.

References

[1] Du, R., Chang, D., Bhunia, A. K., Xie, J., Ma, Z., Song, Y. Z., & Guo, J. (2020). Fine-Grained Visual Classification via Progressive Multi-granularity Training of Jigsaw Patches. In *Computer Vision – ECCV 2020* (pp. 153–168). Springer, Cham.

[2] Yan, Z., Zhang, H., Piramuthu, R., Jagadeesh, V., DeCoste, D., Di, W., & Yu, Y. (2015). HD-CNN: Hierarchical deep convolutional neural network for large-scale visual recognition. In *Proceedings of the IEEE International Conference on Computer Vision* (pp. 2740–2748).

[3] Hu, J., Shen, L., & Sun, G. (2018). Squeeze-and-excitation networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* (pp. 7132–7141).

[4] Woo, S., Park, J., Lee, J. Y., & Kweon, I. S. (2018). CBAM: Convolutional block attention module. In *Proceedings of the European Conference on Computer Vision* (pp. 3–19).

[5] Zheng, H., Fu, J., Mei, T., & Luo, J. (2017). Learning multi-attention convolutional neural network for fine-grained image recognition. In *Proceedings of the IEEE International Conference on Computer Vision* (pp. 5209–5217).

[6] Fu, J., Zheng, H., & Mei, T. (2017). Look closer to see better: Recurrent attention convolutional neural network for fine-grained image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* (pp. 4438–4446).

[7] He, J., Chen, J. N., Liu, S., Kortylewski, A., Yang, C., Bai, Y., & Wang, C. (2021). TransFG: A Transformer Architecture for Fine-grained Recognition. In *Proceedings of the AAAI Conference on Artificial Intelligence* (Vol. 35, No. 1, pp. 852–860).

[8] Zhang, N., Donahue, J., Girshick, R., & Darrell, T. (2014). Part-based R-CNNs for fine-grained category detection. In *European Conference on Computer Vision* (pp. 834–849). Springer, Cham.

[9] Lin, T. Y., RoyChowdhury, A., & Maji, S. (2015). Bilinear CNN models for fine-grained visual recognition. In *Proceedings of the IEEE International Conference on Computer Vision* (pp. 1449–1457).

[10] Wang, X., Girshick, R., Gupta, A., & He, K. (2018). Non-local neural networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* (pp. 7794–7803).

[11] Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J., & Houlsby, N. (2021). An Image Is Worth 16x16 Words: Transformers for Image Recognition at Scale. In *International Conference on Learning Representations*.

[12] Liu, Z., Lin, Y., Cao, Y., Hu, H., Wei, Y., Zhang, Z., ... & Guo, B. (2021). Swin Transformer: Hierarchical Vision Transformer using Shifted Windows. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* (pp. 10012–10022).

[13] Touvron, H., Cord, M., Douze, M., Massa, F., Sablayrolles, A., & Jégou, H. (2021). Training data-efficient image transformers & distillation through attention. In *International Conference on Machine Learning* (pp. 10347–10357). PMLR.

[14] Liu, Z., Mao, H., Wu, C. Y., Feichtenhofer, C., Darrell, T., & Xie, S. (2022). A ConvNet for the 2020s. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 11976–11986).

[15] Wang, W., Xie, E., Li, X., Fan, D. P., Song, K., Liang, D., ... & Shao, L. (2022). PVT v2: Improved Baselines with Pyramid Vision Transformer. *Computational Visual Media*, 8(3), 415–424.

[16] Woo, S., Debnath, S., Hu, R., Chen, X., Liu, Z., Kweon, I. S., & Xie, S. (2023). ConvNeXt V2: Co-Designing and Scaling ConvNets With Masked Autoencoders. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 16133–16142).

[17] Wang, X., Sun, Y., Liu, X., Zou, Z., Wang, L., Wang, K., Liang, X., & Liu, W. (2025). Rethinking Attention Mechanism: Channel Re-attention and Spatial Multi-region Attention for Fine-grained Visual Classification. *Neural Processing Letters*, 57, 43.

[18] Wang, Y., Wang, Y., Ye, Q., Lang, W., & Xu, C. (2025). Strengthen contrastive semantic consistency for fine-grained image classification. *Pattern Analysis and Applications*, 28, 75.

[19] Gao, P., Liang, Y., Song, Y., & Yang, Z. (2025). Cross-Hierarchical Bidirectional Consistency Learning for Fine-Grained Visual Classification. *Information Sciences*, 721, 122562.

[20] Xie, Y., Zeng, Z., Zhang, H., Ding, Y., Wang, Y., Wang, Z., Chen, B., & Liu, H. (2025). Discovering Fine-Grained Visual-Concept Relations by Disentangled Optimal Transport Concept Bottleneck Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 30199–30209).

[21] Park, S., Wang, Z., & Yu, S. X. (2026). Free-Grained Hierarchical Recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

[22] Sikdar, A., Liu, Y., Kedarisetty, S., et al. (2025). Interweaving Insights: High-Order Feature Interaction for Fine-Grained Visual Recognition. *International Journal of Computer Vision*, 133(4), 1755–1779.